

# cancer\_deathrate\_prediction\_OLS\_regression

April 24, 2018

```
In [1]: %matplotlib inline
import pandas as pd
import numpy as np
from sklearn import linear_model
from sklearn import ensemble
from sklearn.model_selection import train_test_split
import matplotlib.pyplot as plt
```

```
In [2]: df = pd.read_csv('cancer_reg.csv')
df.head()
```

```
Out[2]:
```

	avgAnnCount	avgDeathsPerYear	TARGET_deathRate	incidenceRate	medIncome	\
0	1397.0	469	164.9	489.8	61898	
1	173.0	70	161.3	411.6	48127	
2	102.0	50	174.7	349.7	49348	
3	427.0	202	194.8	430.4	44243	
4	57.0	26	144.4	350.1	49955	

  

	popEst2015	povertyPercent	studyPerCap	binmedInc	MedianAge	\
0	260131	11.2	499.748204	(61494.5, 125635]	39.3	
1	43269	18.6	23.111234	(48021.6, 51046.4]	33.0	
2	21026	14.6	47.560164	(48021.6, 51046.4]	45.0	
3	75882	17.1	342.637253	(42724.4, 45201]	42.8	
4	10321	12.5	0.000000	(48021.6, 51046.4]	48.3	

  

	...	PctPrivateCoverageAlone	PctEmpPrivCoverage	PctPublicCoverage	\
0	...	NaN	41.6	32.9	
1	...	53.8	43.6	31.1	
2	...	43.5	34.9	42.1	
3	...	40.3	35.0	45.3	
4	...	43.9	35.1	44.0	

  

	PctPublicCoverageAlone	PctWhite	PctBlack	PctAsian	PctOtherRace	\
0	14.0	81.780529	2.594728	4.821857	1.843479	
1	15.3	89.228509	0.969102	2.246233	3.741352	
2	21.1	90.922190	0.739673	0.465898	2.747358	
3	25.0	91.744686	0.782626	1.161359	1.362643	

```
4                22.7  94.104024  0.270192  0.665830        0.492135
```

```
    PctMarriedHouseholds  BirthRate
0          52.856076    6.118831
1          45.372500    4.333096
2          54.444868    3.729488
3          51.021514    4.603841
4          54.027460    6.796657
```

```
[5 rows x 34 columns]
```

```
In [3]: # missing value handling
```

```
df['PctSomeCol18_24'] = df['PctSomeCol18_24'].fillna(df['PctSomeCol18_24'].mean())
df['PctEmployed16_Over'] = df['PctEmployed16_Over'].fillna(df['PctEmployed16_Over'].mean())
df['PctPrivateCoverageAlone'] = df['PctPrivateCoverageAlone'].fillna(df['PctPrivateCoverageAlone'].mean())
```

```
In [4]: y = df['TARGET_deathRate']
y.head()
```

```
Out[4]: 0    164.9
1    161.3
2    174.7
3    194.8
4    144.4
Name: TARGET_deathRate, dtype: float64
```

```
In [5]: x = df
x = x.drop('TARGET_deathRate', axis=1)
x = x.drop('binnedInc', axis=1) #removed binnedInc field
x.head()
x = pd.get_dummies(x)
x.head()
```

```
Out[5]:   avgAnnCount  avgDeathsPerYear  incidenceRate  medIncome  popEst2015  \
0         1397.0             469           489.8       61898      260131
1         173.0             70           411.6       48127      43269
2         102.0             50           349.7       49348      21026
3         427.0            202           430.4       44243      75882
4          57.0             26           350.1       49955      10321

   povertyPercent  studyPerCap  MedianAge  MedianAgeMale  MedianAgeFemale  \
0          11.2    499.748204       39.3         36.9         41.7
1          18.6    23.111234       33.0         32.2         33.7
2          14.6    47.560164       45.0         44.0         45.8
3          17.1   342.637253       42.8         42.2         43.4
4          12.5     0.000000       48.3         47.8         48.9

   Geography_York County, Pennsylvania  \
0                                     0
```

1	...	0
2	...	0
3	...	0
4	...	0

  

	Geography_York County, South Carolina	Geography_York County, Virginia \
0	0	0
1	0	0
2	0	0
3	0	0
4	0	0

  

	Geography_Young County, Texas	Geography_Yuba County, California \
0	0	0
1	0	0
2	0	0
3	0	0
4	0	0

  

	Geography_Yukon-Koyukuk Census Area, Alaska \
0	0
1	0
2	0
3	0
4	0

  

	Geography_Yuma County, Arizona	Geography_Yuma County, Colorado \
0	0	0
1	0	0
2	0	0
3	0	0
4	0	0

  

	Geography_Zapata County, Texas	Geography_Zavala County, Texas
0	0	0
1	0	0
2	0	0
3	0	0
4	0	0

[5 rows x 3078 columns]

In [6]: train\_x,test\_x,train\_y,test\_y = train\_test\_split(x,y)

In [7]: lr = linear\_model.LinearRegression()  
lr.fit(train\_x,train\_y)

Out[7]: LinearRegression(copy\_X=True, fit\_intercept=True, n\_jobs=1, normalize=False)

```
In [8]: lr.intercept_
```

```
Out[8]: 166.75087262224966
```

```
In [9]: lr.coef_
```

```
Out[9]: array([-2.86465051e-03,  1.60135840e-02,  1.92088818e-01, ...,  
              1.48800406e+01,  0.00000000e+00, -4.57450314e+01])
```

```
In [10]: print('new data score: ',lr.score(test_x,test_y), 'same data score: ',lr.score(train_x,  
( 'new data score: ', 0.4270123712613325, 'same data score: ', 1.0)
```